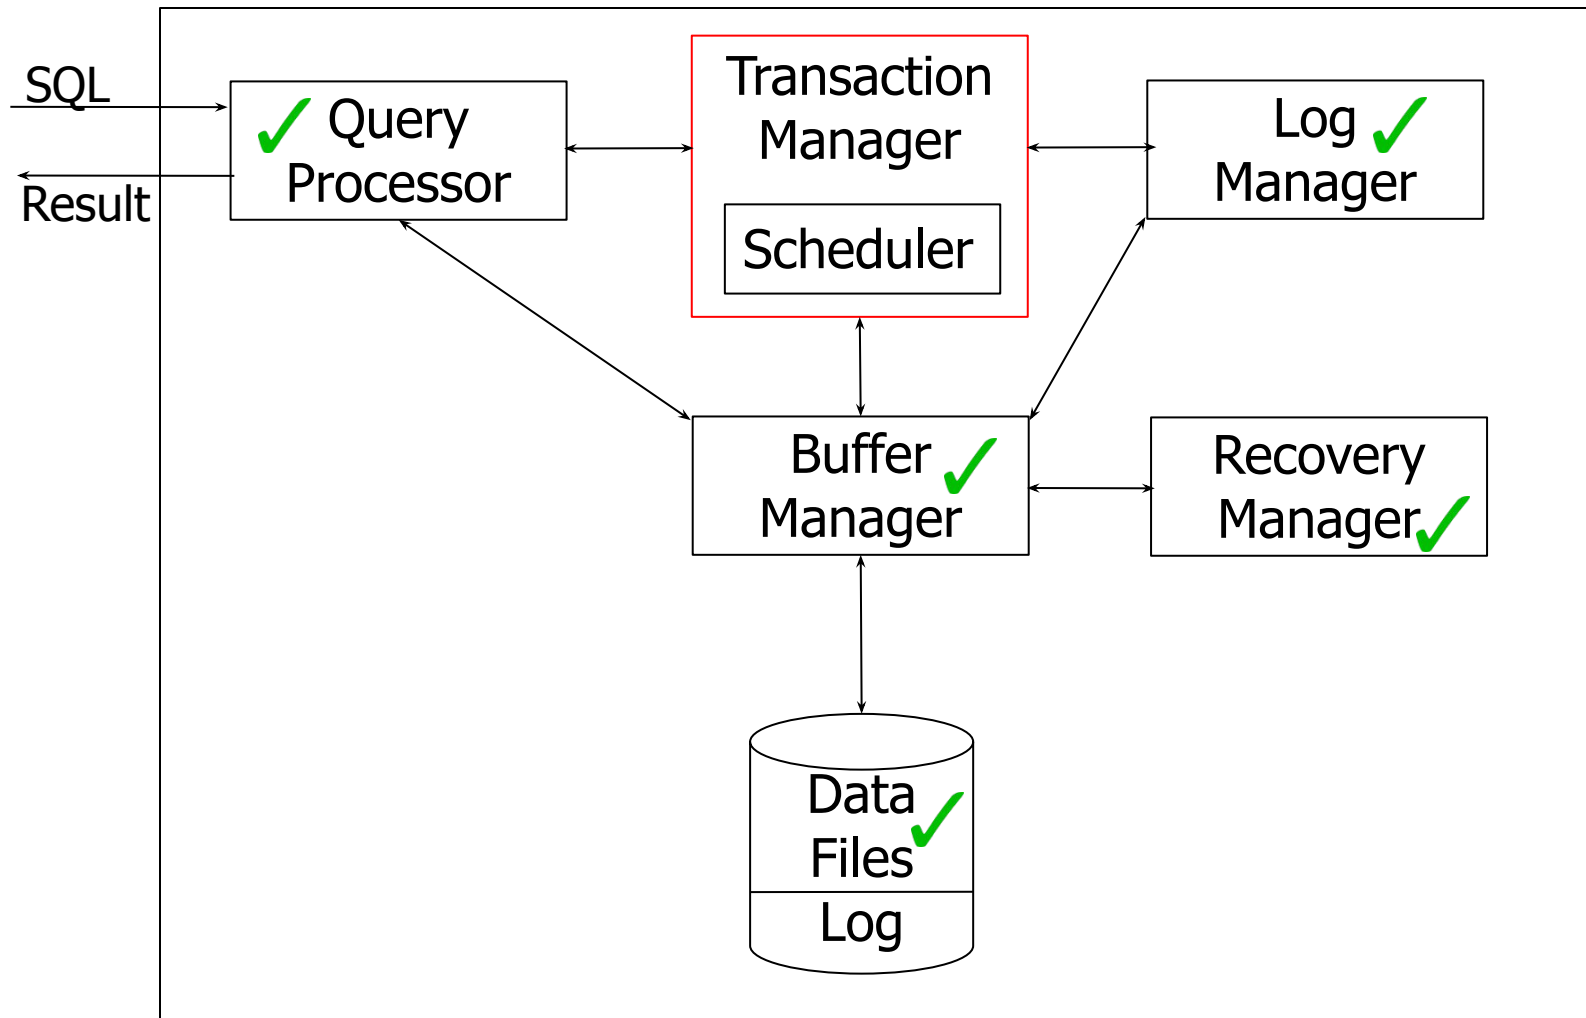


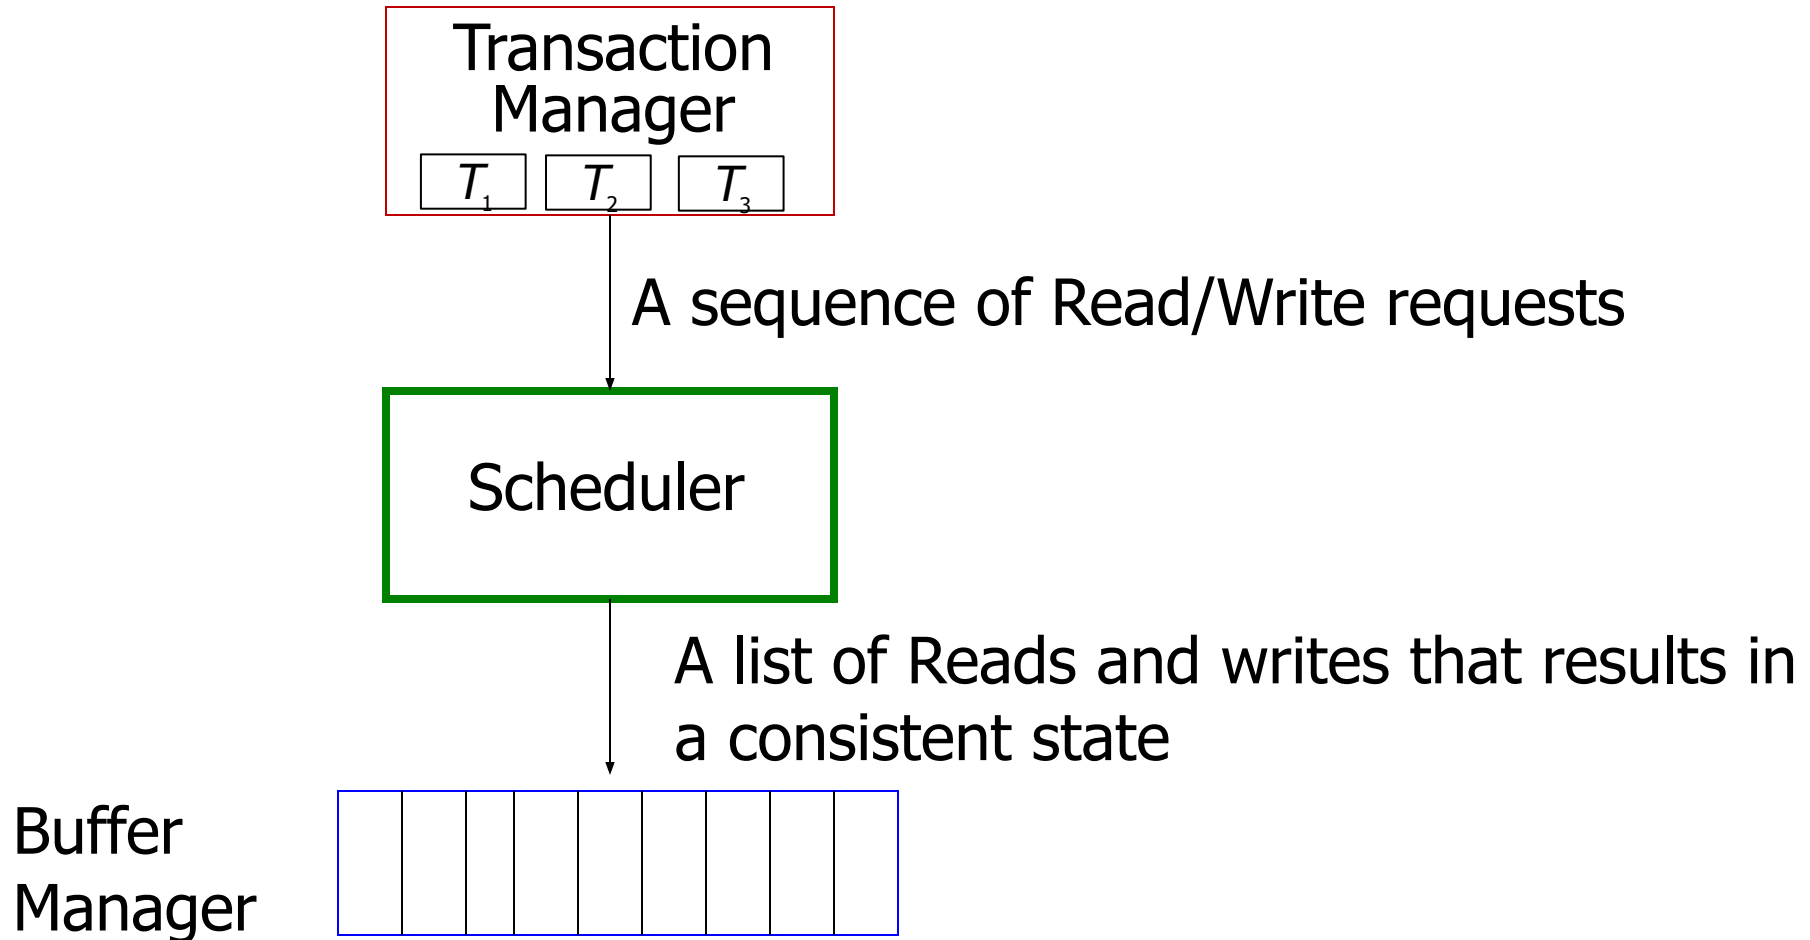
# Lecture 9

## Transactions

Dr. Caroline Sabty  
caroline.sabty@giu-uni.de  
Faculty of Informatics and Computer Science  
German International University in Cairo

These slides are the slides of the course of CSEN 604 Data Bases II taught by Dr. Wael Abouelsaadat that are based on the book of Database Systems; the Complete Book






- Historical note:
  - Turing Award for Transaction concept
  - Jim Gray (1998)

- Interesting reading (optional):

Transaction Concept: Virtues and Limitations by Jim Gray  
<http://www.hpl.hp.com/techreports/tandem/TR-81.3.pdf>

Bank database: 3 Accounts



Property:  $A + B + C = 1500$

Money should not leave/enter

## Transaction T1: Transfer 100 from A to B


$A = 500, B = 500, C = 500$   $\longrightarrow$

Read (A, t)  
 $t = t - 100$   
Write (A, t)  
Read (B, t)  
 $t = t + 100$   
Write (B, t)

$A = 400, B = 600, C = 500$   $\longrightarrow$

## Transaction T2: Transfer 100 from A to C

$A = 500, B = 500, C = 500$



Read (A, s)

$s = s - 100$


Write (A, s)

Read (C, s)

$s = s + 100$

Write (C, s)

$A = 400, B = 500, C = 600$



Transaction T1

Transaction T2

A

B

C

Read (A, t)

500

500

500

$t = t - 100$

Write (A, t)

400

500

500

Read (B, t)

$t = t + 100$

Write (B, t)

400

600

500

Read (A, s)

$s = s - 100$

Write (A, s)

Read (C, s)

$s = s + 100$

Write (C, s)

300

600

500

300

600

600

$$300 + 600 + 600 = 1500$$



Transaction T1

Transaction T2

A

B

C

Read (A, t)

$t = t - 100$

Write (A, t)

Read (A, s)

$s = s - 100$

Write (A, s)

Read (B, t)

$t = t + 100$

Write (B, t)

Read (C, s)

$s = s + 100$

Write (C, s)

500

500

500

400

500

500

300

500

500

300

600

500

300

600

600

$300 + 600 + 600 = 1500$

Transaction T1

Transaction T2

A

B

C

Read (A, t)

$t = t - 100$

Read (A, s)

$s = s - 100$

Write (A, s)

Write (A, t)

Read (B, t)

$t = t + 100$

Write (B, t)

Read (C, s)

$s = s + 100$

Write (C, s)

500

500

500

400

500

500

400

500

500

400

600

500


400

600

600

$400 + 600 + 600 = 1600$

- **Schedule:**
  - The exact sequence of (relevant) actions of one or more transactions
  - The database engine scheduler task is to produce interleaving of transactions steps that result in a consistent database state



- Which schedules are “correct”?
  - Mathematical characterization
- How to build a system that allows only “correct” schedules?
  - Efficient procedure to enforce correctness

# Serial Schedule

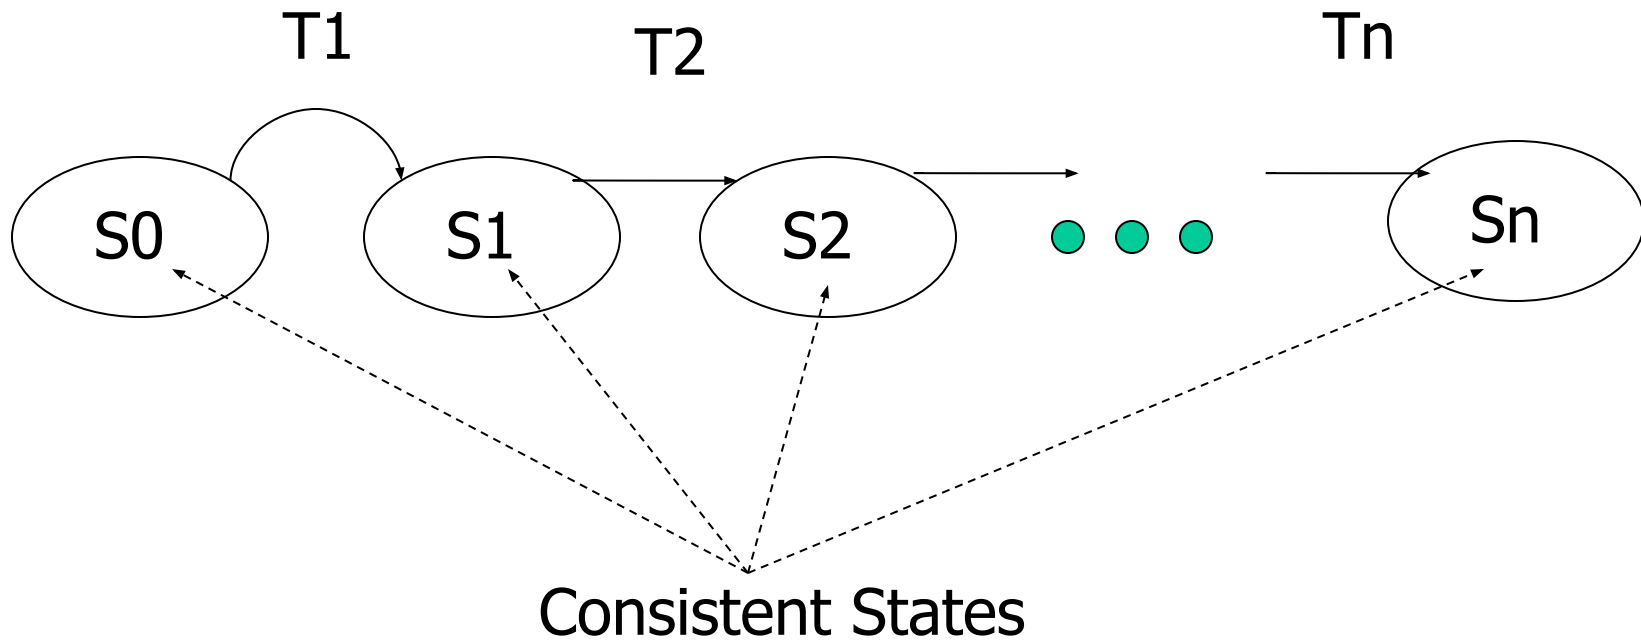
		A	B	C
	Read (A, t)	500	500	500
	$t = t - 100$			
T1	Write (A, t)			
	Read (B, t)			
	$t = t + 100$			
	Write (B, t)	400	600	500
	Read (A, s)			
	$s = s - 100$			
	Write (A, s)			
T2	Read (C, s)			
	$s = s + 100$			
	Write (C, s)	300	600	600

$$300 + 600 + 600 = 1500$$


# Serial Schedule

		A	B	C
T2	Read (A, s)	500	500	500
	$s = s - 100$			
	Write (A, s)			
	Read (C, s)			
	$s = s + 100$			
	Write (C, s)	400	500	600
T1	Read (A, t)			
	$t = t - 100$			
	Write (A, t)			
	Read (B, t)			
	$t = t + 100$			
	Write (B, t)	300	600	600

$$300 + 600 + 600 = 1500$$



- If any action of transaction  $T_1$  precedes any action of  $T_2$ , then all action of  $T_1$  precede all action of  $T_2$
- The correctness principle tells us that every serial schedule will preserve consistency of the database state



- What's the problem with a Serial Schedule?

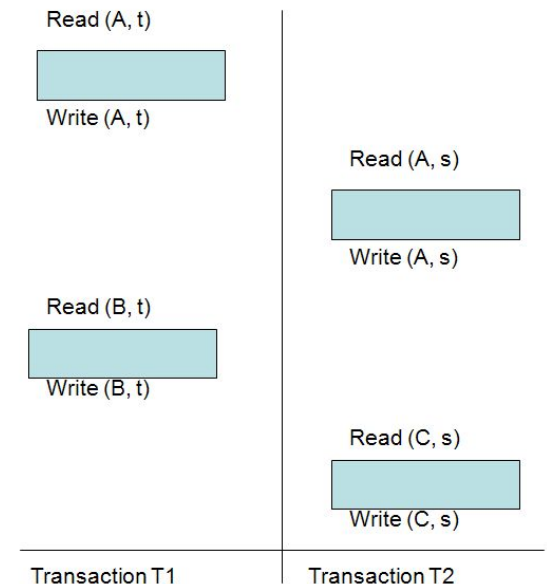


- A serial schedule is the gold standard as it guarantees that the database is in a consistent state. However, it is too slow to use since it means running one transaction at a time and thus no support for concurrency

- A schedule is called *serializable* if its final effect is the same as that of a *serial schedule*
- Serializability: schedule is fine and does not result in inconsistent database
  - Since serial schedules are fine
- Non-serializable schedules are unlikely to result in consistent databases
- Scheduler ensures serializability

- Not possible to look at all  $n!$  serial schedules to check if the effect is the same
  - Instead we ensure serializability by allowing or not allowing certain schedules

- Weaker notion of serializability
- Depends only on reads and writes
- Which steps can be interleaved and which cannot



- Recall from OS course:
  - Multitasking
  - context switch

T1	T2
read(A) A = A - 50 write(A)  <b>read(B)</b> <b>B = B + 50</b> <b>write(B)</b>	<b>read(A)</b> <b>tmp = A * 0.1</b> <b>A = A - tmp</b> <b>write(A)</b>  read(B) B = B + tmp write(B)
Effect:	<div> <div><u>Before</u></div> <div><u>After</u></div> </div>
A	100 45
B	50 105

T1	T2
read(A) A = A - 50 write(A)	read(A) tmp = A*0.1 A = A - tmp write(A)
read(B) B=B+50 write(B)	read(B) B = B+ tmp write(B)

T1	T2
read(A) A = A - 50 write(A)	<b>read(A)</b> <b>tmp = A*0.1</b> <b>A = A - tmp</b> <b>write(A)</b>
<b>read(B)</b> <b>B=B+50</b> <b>write(B)</b>	read(B) B = B+ tmp write(B)

Effect:	<u>Before</u>	<u>After</u>	
A	100	45	==
B	50	105	

Effect:	<u>Before</u>	<u>After</u>
A	100	45
B	50	105

T1	T2
read(A) A = A - 50 write(A)	read(A) tmp = A*0.1 A = A - tmp write(A)
read(B) B=B+50 write(B)	read(B) B = B+ tmp write(B)

T1	T2
read(A) A = A - 50 write(A)	read(A) tmp = A*0.1 A = A - tmp
<b>read(B)</b> B=B+50 write(B)	<b>write(A)</b> read(B) B = B+ tmp write(B)

Effect:	<u>Before</u>	<u>After</u>
A	100	45
B	50	105

==

Effect:	<u>Before</u>	<u>After</u>
A	100	45
B	50	105



T1	T2
read(A) $A = A - 50$ write(A)	read(A) $tmp = A * 0.1$ $A = A - tmp$ write(A)
read(B) $B = B + 50$ write(B)	read(B) $B = B + tmp$ write(B)

T1	T2
read(A) $A = A - 50$ write(A)	read(A) $tmp = A * 0.1$ $A = A - tmp$ write(A)
read(B) $B = B + 50$ <b>write(B)</b>	<b>read(B)</b> $B = B + tmp$ write(B)

Effect:    Before    After

A	100	45
B	50	105

! ==

Effect:    Before    After

A	100	45
B	50	55

T1	T2
read(A) A = A - 50 write(A)	read(A) tmp = A*0.1 A = A - tmp write(A)
read(B) B=B+50 write(B)	read(B) B = B+ tmp write(B)

T1	T2
read(A) A = A - 50 write(A)	read(A) tmp = A*0.1 A = A - tmp
read(B) <b>B=B+50</b> write(B)	<b>write(A)</b>  read(B) B = B+ tmp write(B)

Effect:	<u>Before</u>	<u>After</u>	
A	100	45	==
B	50	105	

Effect:	<u>Before</u>	<u>After</u>
A	100	45
B	50	105

- A conflict serializable schedule produce the same result as one of the serial schedules. It is constructed by considering non-conflicting steps.

$r_T(X)$

Transaction T reads X

$w_T(X)$

Transaction T writes X

- X could be any component of a database:
  - Attribute of a tuple
  - Tuple
  - Block in which a tuple resides
  - A relation
  - ...

- Two Reads
  - E.g.,  $r_i(X); r_i(Y)$
- Read and write of different database element
  - E.g.,  $r_i(X); w_j(Y)$
- Two writes of different database elements
  - E.g.,  $w_i(X); w_j(Y)$

- Two actions of the same transaction
  - E.g.,  $r_i(X); w_i(Y)$
- Two writes of the same database element
  - E.g.,  $w_i(X); w_j(X)$
- A read and a write of the same database element
  - E.g.,  $r_i(X); w_j(X)$

- Conflict-equivalent schedules:
  - If  $S$  can be transformed into  $S'$  through a series of swaps,  $S$  and  $S'$  are called *conflict-equivalent*
  - *conflict-equivalent guarantees same final effect on the database*
- A schedule  $S$  is **conflict-serializable** if it is conflict-equivalent to a serial schedule




- Commercial systems generally support *conflict-serializability*
- Turn a given schedule to a serial one by make as many nonconflicting swaps as we wish

- Given a schedule, determine if it is conflict-serializable
- Construct a *precedence-graph* over the transactions
  - A directed edge from T1 and T2, if they have conflicting instructions, and T1's conflicting instruction comes first
- If there is a cycle in the graph → not conflict-serializable
  - Can be checked in at most  $O(n+e)$  time, where  $n$  is the number of vertices, and  $e$  is the number of edges
- If there is none → conflict-serializable

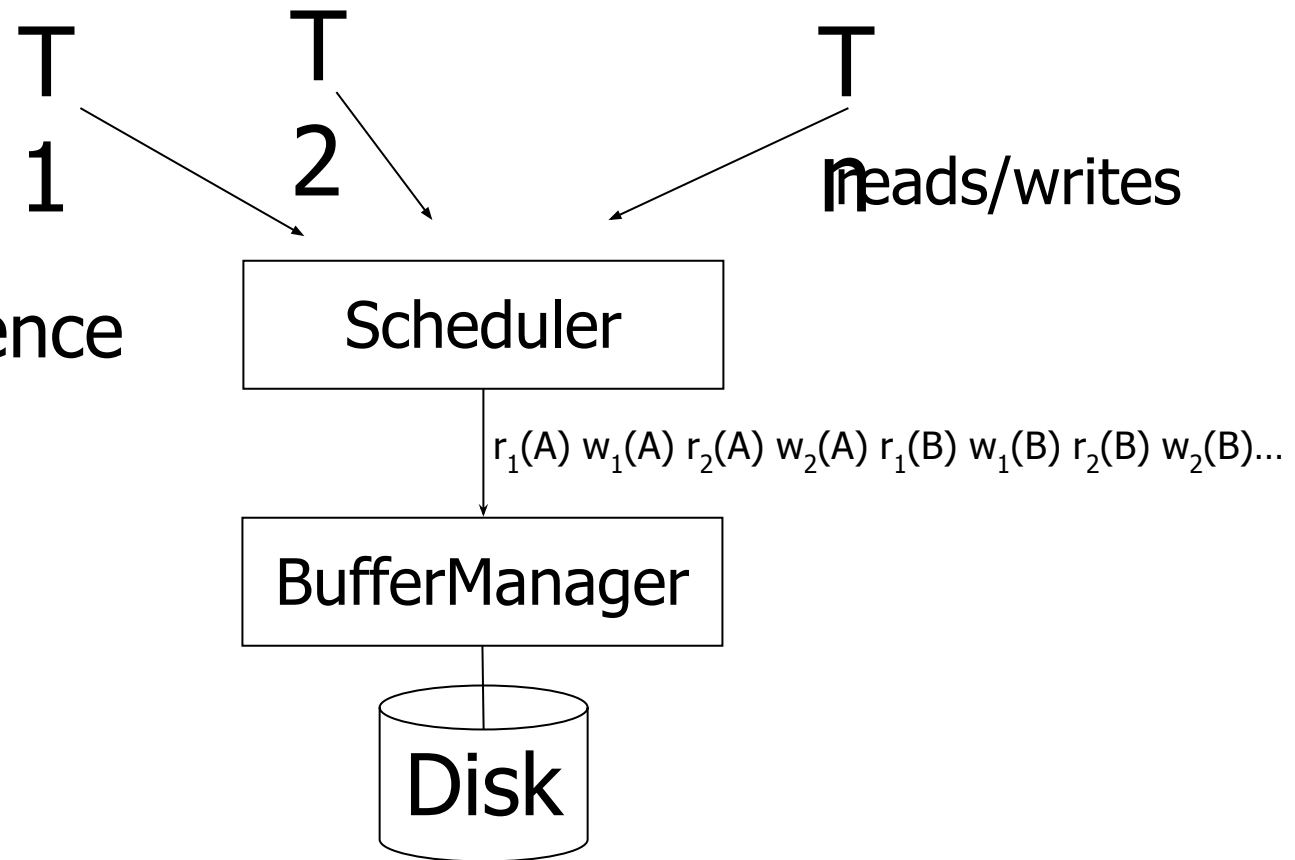
- Precedence graph for schedule S:

- Nodes: Transactions in S
- Edges:  $T_i \rightarrow T_j$  whenever
  - S: ...  $r_i(X)$  ...  $w_i(X)$  ...
  - S: ...  $w_i(X)$  ...  $w_j(X)$  ...
  - S: ...  $r_i(X)$  ...  $w_j(X)$  ...




Note: not necessarily consecutive

Strategy:  
Prevent precedence  
graph cycles




# Graph Theory 101

## Directed Graph:




# Graph Theory 101

## Directed Graph:




# Graph Theory 101

## Directed Graph:



# Graph Theory 101

## Directed Graph:






# Graph Theory 101

Acyclic Graph: A graph with no cycles

# Graph Theory 101

Acyclic  
Graph:



- $T_i \rightarrow T_j$  whenever:
  - There is an action of  $T_i$  that occurs before a conflicting action of  $T_j$ .

Conflicting Steps

$r_i(X); w_i(Y)$
$w_i(X); w_j(X)$
$r_i(X); w_j(X)$

- $T_i \rightarrow T_j$  whenever:
  - There is an action of  $T_i$  that occurs before a conflicting action of  $T_j$ .


$S_1: r_2(A); r_1(B); w_2(A); r_3(A); w_1(B); w_3(A); r_2(B); w_2(B);$

Conflicting Steps

$r_i(X); w_i(Y)$
$w_i(X); w_j(X)$
$r_i(X); w_j(X)$

- $T_i \rightarrow T_j$  whenever:
  - There is an action of  $T_i$  that occurs before a conflicting action of  $T_j$ .

$S_1: r_2(A); r_1(B); w_2(A); r_3(A); w_1(B); w_3(A); r_2(B); w_2(B);$


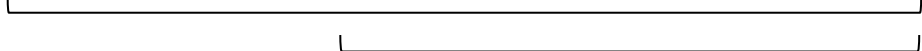


Conflicting Steps

$r_i(X);$	$w_i(Y)$
$w_i(X);$	$w_j(X)$
$r_i(X);$	$w_j(X)$

- $T_i \rightarrow T_j$  whenever:
  - There is an action of  $T_i$  that occurs before a conflicting action of  $T_j$ .

$S_1: r_2(A); r_1(B); w_2(A); r_3(A); w_1(B); w_3(A); r_2(B); w_2(B);$




Conflicting Steps

$r_i(X); w_i(Y)$
$w_i(X); w_j(X)$
$r_i(X); w_j(X)$

- $T_i \rightarrow T_j$  whenever:
  - There is an action of  $T_i$  that occurs before a conflicting action of  $T_j$ .

$S_1: r_2(A); r_1(B); w_2(A); r_3(A); w_1(B); w_3(A); r_2(B); w_2(B);$



Conflicting Steps

$r_i(X); w_i(Y)$
$w_i(X); w_j(X)$
$r_i(X); w_j(X)$

- $T_i \rightarrow T_j$  whenever:
  - There is an action of  $T_i$  that occurs before a conflicting action of  $T_j$ .

$S_2: r_2(A); r_1(B); w_2(A); r_2(B); r_3(A); w_1(B); w_3(A); w_2(B);$


Conflicting Steps

$r_i(X); w_i(Y)$
$w_i(X); w_j(X)$
$r_i(X); w_j(X)$



- $T_i \rightarrow T_j$  whenever:
  - There is an action of  $T_i$  that occurs before a conflicting action of  $T_j$ .

$S_2: r_2(A); r_1(B); w_2(A); r_2(B); r_3(A); w_1(B); w_3(A); w_2(B);$




Conflicting Steps

$r_i(X); w_i(Y)$
$w_i(X); w_j(X)$
$r_i(X); w_j(X)$

- $T_i \rightarrow T_j$  whenever:
  - There is an action of  $T_i$  that occurs before a conflicting action of  $T_j$ .

$S_2: r_2(A); r_1(B); w_2(A); r_2(B); r_3(A); w_1(B); w_3(A); w_2(B);$




Conflicting Steps

$r_i(X); w_i(Y)$
$w_i(X); w_j(X)$
$r_i(X); w_j(X)$

- $T_i \rightarrow T_j$  whenever:
  - There is an action of  $T_i$  that occurs before a conflicting action of  $T_j$ .

$S_2: r_2(A); r_1(B); w_2(A); r_2(B); r_3(A); w_1(B); w_3(A); w_2(B);$




Conflicting Steps

$r_i(X); w_i(Y)$
$w_i(X); w_j(X)$
$r_i(X); w_j(X)$

# Precedence Graph – Example 3

$T_1$	$T_2$	$T_3$	$T_4$	$T_5$
read(Y) read(Z)	read(X)		read(V) read(W) read(W)	
	read(Y) write(Y)	write(Z)		
read(U)		read(Y) write(Y) read(Z) write(Z)		
read(U) write(U)				



Conflicting Steps

$r_i(X); w_i(Y)$   
 $w_i(X); w_j(X)$   
 $r_i(X); w_j(X)$

$r_2(X); r_1(Y); r_1(Z); r_5(V); r_5(W); r_5(W); r_2(Y); w_2(Y); w_3(Z); r_1(U); r_4(Y); w_4(Y); r_4(Z); w_4(Z); r_1(U); w_1(U)$

- Two schedules are *conflict-equivalent* if they can be turned one into the other by a sequence of nonconflicting swaps of adjacent actions
- A schedule is *conflict-serializable* if it is conflict-equivalent to a serial schedule

Thank you !